

The Google File System & A Comparison of Approaches to Large-Scale Data Analysis

Nicholas D. Bradford

10.20.16

The Google File System – Main Idea

- Developers at Google recognized that their data storage system was no longer satisfying their growing needs as a company.
- Google File System was designed and implemented to be a scalable file system.
- The system was intended to be running on inexpensive hardware and the delivery of high aggregate performance.
- GFS was designed to work using a sprawling system of a thousand smaller, faster processing servers that could be repaired and rerouted with ease.
- GFS was also made to work to handle larger files, append data instead of rewriting it, and allow for simultaneous access of various data and programs.

The Google File System - Implementation

- The GFS is physically designed using inexpensive commodity components that commonly fail and must be replaced or waited on to be repaired.
- GFS stores larger than traditional files ex. Multi-GB files
- GFS handles two types of access traffic, large streaming read which includes reading MBs and small random reads that read small KB bits of data.
- Workloads that support large writes and appending data files
- Allow for multiple clients to simultaneously append data
- Bandwidth > Latency

The Google File System - Analysis

- A useful, cheap, and maneuverable system that fulfills Google's need
- A more powerful system than previous versions due to ability to store and access larger bulk files
- Design concepts that allow for the appending of data and therefore less processing time when updating the system
- Overall: **B+** System due to potential faults of using less reliable but cheaper machines

A Comparison of Approaches to Large-Scale Data Analysis – Main Idea

- A group of professors and researchers from various Universities and Companies ran experiments to evaluate both the MapReduce and the Database Management System paradigms of data-analysis.

A Comparison of Approaches to Large-Scale Data Analysis – Implementation

- The experiment was run on an open source Hadoop MR server and two parallel SQL DBMSs (Vertica and a major retail vendor)
- Benchmarks were used to establish a baseline test for both machines.
- Performance measured for parallelism against a cluster of 100 nodes

A Comparison of Approaches to Large-Scale Data Analysis – Analysis

Google File System/A Comparison of Approaches to Large-Scale Data Analysis – Comparison

- GFS is a MapReduce system that works well for Google
- DBMS systems are more expensive and harder to implement but in the long run more successful

Stonbraker Talk – Main Ideas

Advantages/Disadvantages (In Context)

Advantages

Disadvantages